

A stylized blue profile of a human head is on the left, with a glowing blue audio waveform extending from the mouth area across the right side of the image. The background is a gradient of dark blue to light blue.

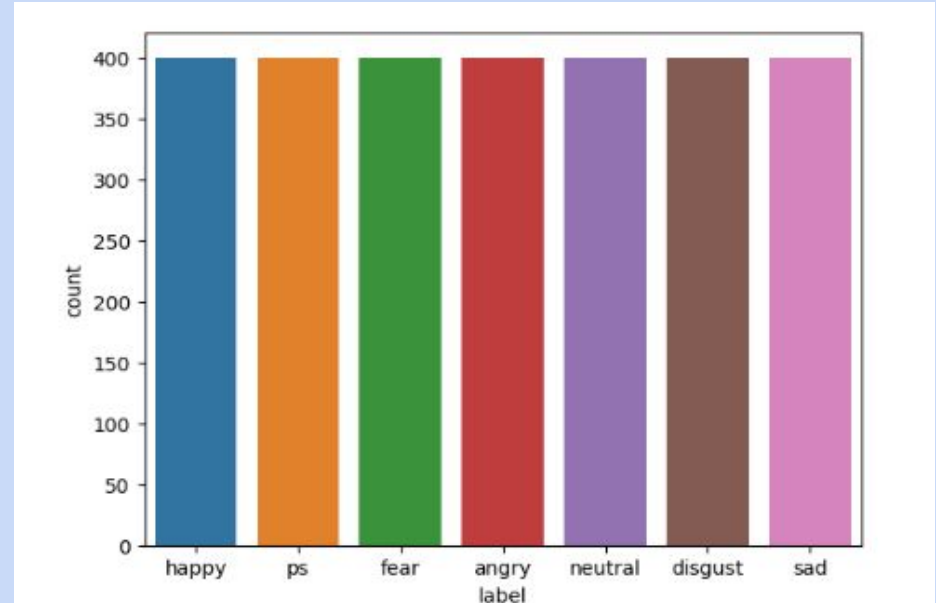
EmoSonix: Audio Emotion Detection System

Presented by

- Akib Rashed (20301220)
- Mashfia Zaman Toa (20301229)
- Arham Ahmed Jobayer (20301216)
- Md Shamiul Islam Khan (20301235)

Previously Done work on Dataset


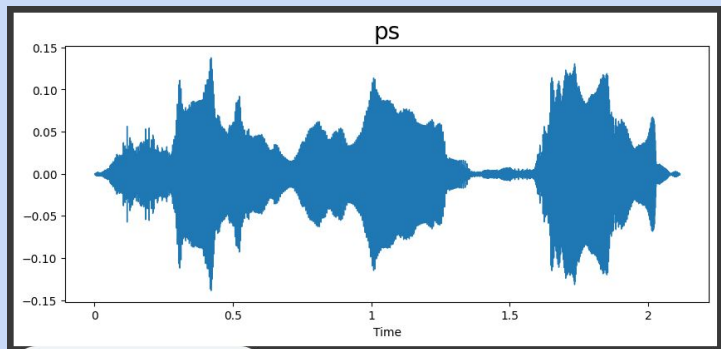
- Loaded the data and separated path and label
- Balanced dataset with 400 data point in each Category




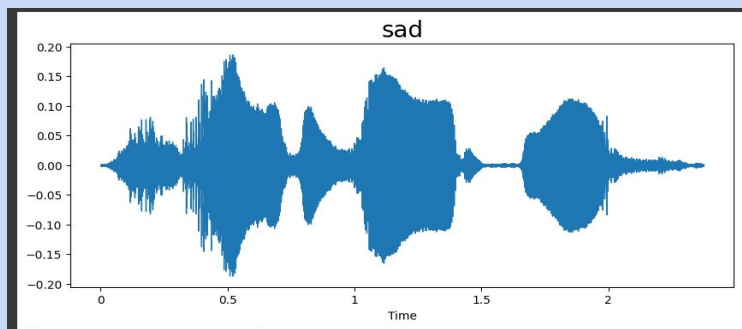


Data Wrangling

- Waveplot
 - Visualization
 - Comparison
 - Transient Detection
 - Noise and Distortion Analysis

```
#waveform of data ar_jonno
def waveplot(data, sr, emotion):
    plt.figure(figsize=(10,4))
    plt.title(emotion, size=20)
    librosa.display.waveshow(data, sr=sr)
    plt.show()
```











Spectrogram


- Frequency Information
- Visualization of Features
- Better Discrimination
- Pattern Recognition
- Feature Extraction

```
def spectrogram(data, sr, emotion):  
    x = librosa.stft(data)  
    xdb = librosa.amplitude_to_db(abs(x))  
    plt.figure(figsize=(11,4))  
    plt.title(emotion, size=20)  
    librosa.display.specshow(xdb, sr=sr, x_axis='time', y_axis='hz')  
    plt.colorbar()
```





Frequency Distribution based on Spectrogram



Feature Extraction

- To convert audio file into numeric format
- Mel-Frequency Cepstral Coefficients (MFCC)

```
array([-3.34945923e+02,  3.70755119e+01, -7.90456820e+00, -1.83790684e+00,  
       -1.28784685e+01,  6.48743868e+00,  2.52078199e+00, -1.22021313e+01,  
       -6.53204012e+00,  1.64398634e+00, -1.77280693e+01,  9.59110737e+00,  
       -9.33374500e+00,  7.03533840e+00, -2.34031916e+00, -5.02835369e+00,  
       -1.51206696e+00,  4.88031101e+00, -2.85840440e+00,  2.46732163e+00,  
       -2.44669652e+00, -1.70657969e+00, -7.33406544e+00, -1.87550449e+00,  
       -2.89007664e+00,  2.28449538e-01,  1.87676847e+00,  7.56279325e+00,  
       1.12165041e+01,  1.49963121e+01,  1.19341040e+01,  1.10935335e+01,  
       4.09022570e+00,  3.95253038e+00,  9.11194146e-01,  4.24853182e+00,  
       3.31549644e+00,  4.12816429e+00, -2.26341224e+00, -1.07533288e+00,  
       -1.64336324e+00,  3.98158669e+00, -1.11243165e+00, -8.54802370e-01,  
       -4.10659170e+00, -8.08297634e-01, -2.51135379e-01,  2.69858575e+00,  
       -1.66069591e+00, -1.15179658e+00,  1.13143539e+00,  1.59848392e+00,  
       -1.75673831e+00, -6.32186159e-02, -1.91152751e+00,  1.58386374e+00,  
       6.33633137e-01,  1.27641773e+00,  7.73862720e-01,  3.91387486e+00],  
      dtype=float32)
```


Normalizing

- Improved Convergence
- Better Performance
- Prevention of Numerical Instabilities

```
] X_mfcc = df['speech'].apply(lambda x: extract_mfcc(x))
```

```
] X_mfcc
```

```
0      [-334.94592, 37.075512, -7.904568, -1.8379068, ...  
1      [-350.88992, 74.09218, -2.3219209, 7.8491626, ...  
2      [-322.57584, 47.85948, -25.91174, 4.019515, -2...  
3      [-289.10486, 78.74502, 2.2043347, -5.4672093, ...  
4      [-325.0482, 52.04364, -21.359823, 3.8501308, -...  
      ...  
2795     [-331.03552, 48.584187, -25.937048, -4.8479657...  
2796     [-340.9328, 93.7364, -26.861814, 0.34102386, -...  
2797     [-354.5709, 64.17273, -31.288214, 14.938522, -...  
2798     [-305.5626, 70.34936, -42.098427, -7.0920105, ...  
2799     [-389.8703, 94.92119, -6.0177207, -5.5790243, ...  
Name: speech, Length: 2800, dtype: object
```

Encoding

- One Hot Encoded into an array

```
#Output Column  
from sklearn.preprocessing import OneHotEncoder  
enc = OneHotEncoder()  
y = enc.fit_transform(df[['label']])
```

```
y = y.toarray()
```

y

```
array([[1., 0., 0., ..., 0., 0., 0.],  
       [1., 0., 0., ..., 0., 0., 0.],  
       [1., 0., 0., ..., 0., 0., 0.],  
       ...,  
       [0., 0., 0., ..., 0., 1., 0.],  
       [0., 0., 0., ..., 0., 1., 0.],  
       [0., 0., 0., ..., 0., 1., 0.]])
```

LSTM on the encoded dataset

- Activation
 - Relu
 - Softmx
- Optimizer
 - Adam
- 50 epoch

```
# Train the model
history = model.fit(X, y, validation_split=0.2, epochs=50, batch_size=64)

Epoch 22/50
35/35 [=====] - 0s 9ms/step - loss: 0.1012 - accuracy: 0.9696 - val_loss: 5.2422 - val_accuracy: 0.2964
Epoch 23/50
35/35 [=====] - 0s 9ms/step - loss: 0.0507 - accuracy: 0.9871 - val_loss: 5.5924 - val_accuracy: 0.2500
Epoch 24/50
35/35 [=====] - 0s 9ms/step - loss: 0.0712 - accuracy: 0.9799 - val_loss: 6.6559 - val_accuracy: 0.0929
Epoch 25/50
35/35 [=====] - 0s 9ms/step - loss: 0.2671 - accuracy: 0.9201 - val_loss: 5.9003 - val_accuracy: 0.1536
Epoch 26/50
35/35 [=====] - 0s 9ms/step - loss: 0.1314 - accuracy: 0.9652 - val_loss: 5.3231 - val_accuracy: 0.2393
Epoch 27/50
35/35 [=====] - 0s 9ms/step - loss: 0.0844 - accuracy: 0.9790 - val_loss: 5.2956 - val_accuracy: 0.2554
Epoch 28/50
35/35 [=====] - 0s 9ms/step - loss: 0.0587 - accuracy: 0.9804 - val_loss: 6.8499 - val_accuracy: 0.1482
Epoch 29/50
35/35 [=====] - 0s 9ms/step - loss: 0.0484 - accuracy: 0.9853 - val_loss: 6.1121 - val_accuracy: 0.2839
Epoch 30/50
35/35 [=====] - 0s 9ms/step - loss: 0.0663 - accuracy: 0.9786 - val_loss: 5.8007 - val_accuracy: 0.2393
Epoch 31/50
35/35 [=====] - 0s 9ms/step - loss: 0.0671 - accuracy: 0.9817 - val_loss: 6.2283 - val_accuracy: 0.2393
Epoch 32/50
35/35 [=====] - 0s 9ms/step - loss: 0.0715 - accuracy: 0.9786 - val_loss: 6.1375 - val_accuracy: 0.2821
Epoch 33/50
35/35 [=====] - 0s 9ms/step - loss: 0.0539 - accuracy: 0.9862 - val_loss: 5.0351 - val_accuracy: 0.3179
Epoch 34/50
35/35 [=====] - 0s 9ms/step - loss: 0.0498 - accuracy: 0.9844 - val_loss: 5.6238 - val_accuracy: 0.3071
Epoch 35/50
35/35 [=====] - 0s 9ms/step - loss: 0.0433 - accuracy: 0.9879 - val_loss: 5.1346 - val_accuracy: 0.3536
Epoch 36/50
35/35 [=====] - 0s 10ms/step - loss: 0.0329 - accuracy: 0.9897 - val_loss: 5.1427 - val_accuracy: 0.3536
Epoch 37/50
```

```
from keras.models import Sequential
from keras.layers import Dense, LSTM, Dropout


model = Sequential([
    LSTM(256, return_sequences=False, input_shape=(60,1)),
    Dropout(0.2),
    Dense(128, activation='relu'),
    Dropout(0.2),
    Dense(64, activation='relu'),
    Dropout(0.2),
    Dense(7, activation='softmax')
])

model.compile(loss='categorical_crossentropy', optimizer='adam', metrics=['accuracy'])
model.summary()
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|---------------------|--------------|---------|
| ===== | | |
| lstm (LSTM) | (None, 256) | 264192 |
| dropout (Dropout) | (None, 256) | 0 |
| dense (Dense) | (None, 128) | 32896 |
| dropout_1 (Dropout) | (None, 128) | 0 |
| dense_1 (Dense) | (None, 64) | 8256 |
| dropout_2 (Dropout) | (None, 64) | 0 |
| dense_2 (Dense) | (None, 7) | 455 |

Results





THANK YOU!